

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report, confusion_matrix
import warnings
import pickle
from scipy import stats
warnings.filterwarnings('ignore')
plt.style.use('fivethirtyeight')
```

```
data=pd.read_csv('/content/Data_Train.csv')
data.head()
```

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Du
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR ? DEL	22:20	01:10 22 Mar	
1	Air India	1/05/2019	Kolkata	Banglore	CCU ? IXR ? BBI ? BLR	05:50	13:15	
2	Jet Airways	9/06/2019	Delhi	Cochin	DEL ? LKO ? BOM ? COK	09:25	04:25 10 Jun	
3	IndiGo	12/05/2019	Kolkata	Banglore	CCU ? NAG ? BLR	18:05	23:30	
4	IndiGo	01/03/2019	Banglore	New Delhi	BLR ? NAG ? DEL	16:50	21:35	

```
data.shape

(10683, 11)
```

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                10683 non-null object
1   Date_of_Journey        10683 non-null object
2   Source                 10683 non-null object
3   Destination            10683 non-null object
4   Route                  10682 non-null object
5   Dep_Time               10683 non-null object
6   Arrival_Time           10683 non-null object
7   Duration               10683 non-null object
8   Total_Stops            10682 non-null object
9   Additional_Info        10683 non-null object
```

```

10 Price          10683 non-null  int64
dtypes: int64(1), object(10)
memory usage: 918.2+ KB

```

```
data.isnull().sum()
```

```

Airline          0
Date_of_Journey  0
Source           0
Destination      0
Route           1
Dep_Time         0
Arrival_Time     0
Duration         0
Total_Stops      1
Additional_Info   0
Price           0
dtype: int64

```

```
category=['Airline','Source','Destination','Additional_Info']
```

```
category
```

```
['Airline', 'Source', 'Destination', 'Additional_Info']
```

```
for i in category:
```

```
    print(i,data[i].unique())
```

```

Airline ['IndiGo' 'Air India' 'Jet Airways' 'SpiceJet' 'Multiple carriers' 'GoAir'
        'Vistara' 'Air Asia' 'Vistara Premium economy' 'Jet Airways Business'
        'Multiple carriers Premium economy' 'Trujet']
Source  ['Bangalore' 'Kolkata' 'Delhi' 'Chennai' 'Mumbai']
Destination ['New Delhi' 'Bangalore' 'Cochin' 'Kolkata' 'Delhi' 'Hyderabad']
Additional_Info ['No info' 'In-flight meal not included' 'No check-in baggage included'
                '1 Short layover' 'No Info' '1 Long layover' 'Change airports'
                'Business class' 'Red-eye flight' '2 Long layover']

```

```
category_cols=data.select_dtypes(include=['object']).columns
```

```
category_cols
```

```

Index(['Airline', 'Date_of_Journey', 'Source', 'Destination', 'Route',
       'Dep_Time', 'Arrival_Time', 'Duration', 'Total_Stops',
       'Additional_Info'],
      dtype='object')

```


```
for column in category_cols:
```

```
    plt.figure(figsize=(20,4))
```

```
    plt.subplot(121)
```

```
    data[column].value_counts().plot(kind='bar')
```

```
    plt.title(column)
```



```
data.Route=data.Route.str.split('->')
data.Route

0          [BLR ? DEL]
1    [CCU ? IXR ? BBI ? BLR]
2    [DEL ? LKO ? BOM ? COK]
3          [CCU ? NAG ? BLR]
4          [BLR ? NAG ? DEL]
...
10678          [CCU ? BLR]
10679          [CCU ? BLR]
10680          [BLR ? DEL]
10681          [BLR ? DEL]
10682    [DEL ? GOI ? BOM ? COK]
Name: Route, Length: 10683, dtype: object

data['City1']=data.Route.str[0]
data['City2']=data.Route.str[1]
data['City3']=data.Route.str[2]
data['City4']=data.Route.str[3]
data['City5']=data.Route.str[4]
data['City6']=data.Route.str[5]

100

data.Date_of_Journey=data.Date_of_Journey.str.split('/')
data.Date_of_Journey

0          [24, 03, 2019]
1          [1, 05, 2019]
2          [9, 06, 2019]
3          [12, 05, 2019]
4          [01, 03, 2019]
...
10678          [9, 04, 2019]
10679          [27, 04, 2019]
10680          [27, 04, 2019]
10681          [01, 03, 2019]
10682          [9, 05, 2019]
Name: Date_of_Journey, Length: 10683, dtype: object
```

```
data['Date']=data.Date_of_Journey.str[0]
data['Month']=data.Date_of_Journey.str[1]
data['Year']=data.Date_of_Journey.str[2]

data.Dep_Time=data.Dep_Time.str.split(':')

data['Dep_Time_Hour']=data.Dep_Time.str[0]
data['Dep_Time_Mins']=data.Dep_Time.str[1]

data.Arrival_Time=data.Arrival_Time.str.split(' ')

data['Arrival_date']=data.Arrival_Time.str[1]
data['Time_of_Arrival']=data.Arrival_Time.str[0]

data['Time_of_Arrival']=data.Time_of_Arrival.str.split(':')

data['Arrival_Time_Hour']=data.Time_of_Arrival.str[0]
data['Arrival_Time_Mins']=data.Time_of_Arrival.str[1]

data['Time_of_Arrival']=data.Time_of_Arrival.str.split(':')

data.Duration=data.Duration.str.split(' ')

data['Travel_Hours']=data.Duration.str[0]
data['Travel_Hours']=data['Travel_Hours'].str.split('h')
data['Travel_Hours']=data['Travel_Hours'].str[0]
data.Travel_Hours=data.Travel_Hours
data['Travel_Mins']=data.Duration.str[1]

data.Travel_Mins=data.Travel_Mins.str.split('m')
data.Travel_Mins=data.Travel_Mins.str[0]

data.Total_Stops.replace('non_stop',0,inplace=True)
data.Total_Stops=data.Total_Stops.str.split(' ')
data.Total_Stops=data.Total_Stops.str[0]

data.Additional_Info.unique()

array(['No info', 'In-flight meal not included',
       'No check-in baggage included', '1 Short layover', 'No Info',
       '1 Long layover', 'Change airports', 'Business class',
       'Red-eye flight', '2 Long layover'], dtype=object)

data.Additional_Info.replace('No Info','No info',inplace=True)

data.isnull().sum()

Airline      0
Date_of_Journey  0
Source       0
Destination  0
Route        1
Dep_Time     0
Arrival_Time 0
Duration     0
Total_Stops  1
Additional_Info 0
Price        0
City1        1
City2      10683
City3      10683
City4      10683
City5      10683
City6      10683
Date        0
```

```

Month                0
Year                0
Dep_Time_Hour        0
Dep_Time_Mins        0
Arrival_date         6348
Time_of_Arrival      10683
Arrival_Time_Hour    0
Arrival_Time_Mins    0
Travel_Hours         0
Travel_Mins          1032
dtype: int64

```

```
data.drop(['City4', 'City5', 'City6'], axis=1, inplace=True)
```

```

data.drop(['Date_of_Journey', 'Route', 'Dep_Time', 'Arrival_Time', 'Duration'], axis=1, inplace=True)
data.drop(['Time_of_Arrival'], axis=1, inplace=True)

```

```
data.isnull().sum()
```

```

Airline                0
Source                0
Destination            0
Total_Stops            1
Additional_Info        0
Price                 0
City1                  1
City2                 10683
City3                 10683
Date                   0
Month                  0
Year                   0
Dep_Time_Hour          0
Dep_Time_Mins          0
Arrival_date           6348
Arrival_Time_Hour      0
Arrival_Time_Mins      0
Travel_Hours           0
Travel_Mins            1032
dtype: int64

```

```
data['City3'].fillna('None', inplace=True)
```

```
data['Arrival_date'].fillna(data['Date'], inplace=True)
```

```
data['Travel_Mins'].fillna(0, inplace=True)
```

```
data.isnull().sum()
```

```

Airline                0
Source                0
Destination            0
Total_Stops            1
Additional_Info        0
Price                 0
City1                  1
City2                 10683
City3                  0
Date                   0
Month                  0
Year                   0
Dep_Time_Hour          0
Dep_Time_Mins          0
Arrival_date           0
Arrival_Time_Hour      0
Arrival_Time_Mins      0
Travel_Hours           0
Travel_Mins            0
dtype: int64

```

```

data.Date=data.Date.astype('int64')
data.Month=data.Month.astype('int64')

```

```
data.Year=data.Year.astype('int64')
data.Dep_Time_Hour=data.Dep_Time_Hour.astype('int64')
data.Dep_Time_Hour=data.Dep_Time_Hour.astype('int64')
data.Dep_Time_Mins=data.Dep_Time_Mins.astype('int64')
data.Arrival_date=data.Arrival_date.astype('int64')
data.Arrival_Time_Hour=data.Arrival_Time_Hour.astype('int64')
data.Arrival_Time_Mins=data.Arrival_Time_Mins.astype('int64')
data.Travel_Mins=data.Travel_Mins.astype('int64')
```

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                10683 non-null  object
1   Source                 10683 non-null  object
2   Destination            10683 non-null  object
3   Total_Stops            10682 non-null  object
4   Additional_Info        10683 non-null  object
5   Price                  10683 non-null  int64
6   City1                  10682 non-null  object
7   City2                  0 non-null      float64
8   City3                  10683 non-null  object
9   Date                   10683 non-null  int64
10  Month                  10683 non-null  int64
11  Year                   10683 non-null  int64
12  Dep_Time_Hour          10683 non-null  int64
13  Dep_Time_Mins          10683 non-null  int64
14  Arrival_date           10683 non-null  int64
15  Arrival_Time_Hour      10683 non-null  int64
16  Arrival_Time_Mins      10683 non-null  int64
17  Travel_Hours           10683 non-null  object
18  Travel_Mins            10683 non-null  int64
dtypes: float64(1), int64(10), object(8)
memory usage: 1.5+ MB
```

```
data[data['Travel_Hours']=='5m']
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2	City3	Date	Month	Year	De
6474	Air India	Mumbai	Hyderabad	2	No info	17327	BOM ? GOI ? PNQ ? HYD	NaN	None	6	3	2019	

```
data.drop(index=6474,inplace=True,axis=0)
```

```
data.Travel_Hours=data.Travel_Hours.astype('int64')
```

```
categorical=['Airline','Source','Destination','Additional_Info','City1']
numerical=['Total_Stops','Date','Month','Year','Dep_Time_Hour','Dep_Time_Mins','Arrival_Time_Mins','Travel_Hours','Trave
```

Double-click (or enter) to edit

```
import seaborn as sns
c=1

plt.figure(figsize=(20,45))
for i in categorical:
    plt.subplot(6,3,c)
    sns.displot('Price')
    plt.xticks(rotation=90)
```

```
plt.tight_layout(pad=3.0)
c=c+1

plt.show()
```

```
data.columns
```

```
Index(['Airline', 'Source', 'Destination', 'Total_Stops', 'Additional_Info',
      'Price', 'City1', 'City2', 'City3', 'Date', 'Month', 'Year',
      'Dep_Time_Hour', 'Dep_Time_Mins', 'Arrival_date', 'Arrival_Time_Hour',
      'Arrival_Time_Mins', 'Travel_Hours', 'Travel_Mins'],
      dtype='object')
```

```
import seaborn as sns
c=1
```

```
for i in categorical:
    plt.figure(figsize = (10,20))
```

```
    plt.subplot(6,3,c)
```

```
sns.scatterplot(x=data[i],y=data.Price)
plt.xticks(rotation=90)
#plt.tight_layout(pad=3.0)
c=c+1
plt.show()
```



```
data[data.Price>50000]
data.head()
pd.set_option('display.max_columns',25)
data.head()
```




```
data['Year'].max()
```

2019

```
sns.heatmap(data.corr(),annot=True)
```

<Axes: >



```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10682 entries, 0 to 10682
Data columns (total 19 columns):
 #   Column                Non-Null Count  Dtype
  ...
```

```
---
0  Airline      10682 non-null object
1  Source      10682 non-null object
2  Destination 10682 non-null object
3  Total_Stops 10681 non-null object
4  Additional_Info 10682 non-null object
5  Price       10682 non-null int64
6  City1       10681 non-null object
7  City2       0 non-null float64
8  City3       10682 non-null object
9  Date        10682 non-null int64
10 Month       10682 non-null int64
11 Year        10682 non-null int64
12 Dep_Time_Hour 10682 non-null int64
13 Dep_Time_Mins 10682 non-null int64
14 Arrival_date 10682 non-null int64
15 Arrival_Time_Hour 10682 non-null int64
16 Arrival_Time_Mins 10682 non-null int64
17 Travel_Hours 10682 non-null int64
18 Travel_Mins 10682 non-null int64
dtypes: float64(1), int64(11), object(7)
memory usage: 1.6+ MB
```

data


```

    Airline  Source  Destination  Total_Stops  Additional_Info  Price  City1  City2  City3  Date  Month  Year
0  IndiGo  Bangalore  New Delhi  0  Non-Stop  6007  BLR  DEL  MUM  04  0  2018
c=1

for i in numerical:
    plt.figure(figsize=(10,20))
    plt.subplot(6,3,c)
    sns.scatterplot(x = data[i], y=data.Price)
    plt.xticks(rotation=90)
    #plt.tight_layout(pad=3.0)
    c=c+1
    plt.show()
```

```

import seaborn as sns
sns.boxplot(data['Price'])
```



```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
```

```
data.Airline=le.fit_transform(data.Airline)
data.Source=le.fit_transform(data.Source)
data.Destination=le.fit_transform(data.Destination)
data.Total_Stops=le.fit_transform(data.Total_Stops)
data.City1=le.fit_transform(data.City1)
data.City2=le.fit_transform(data.City2)
data.City3=le.fit_transform(data.City3)
data.Additional_Info=le.fit_transform(data.Additional_Info)
data.head()
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2	City3	Date	Month	Year	Dep_T:
0	3	0	5	4	7	3897	18	0	0	24	3	2019	
1	1	3	0	1	7	7662	84	0	0	1	5	2019	
2	4	2	1	1	7	13882	118	0	0	9	6	2019	
3	3	3	0	0	7	6218	91	0	0	12	5	2019	
4	3	0	5	0	7	13302	29	0	0	1	3	2019	

```
data = data[['Airline', 'Source', 'Destination', 'Date', 'Month', 'Year', 'Dep_Time_Hour', 'Dep_Time_Mins', 'Arrival_date', 'Arri
```

```
data.head()
```

	Airline	Source	Destination	Date	Month	Year	Dep_Time_Hour	Dep_Time_Mins	Arrival_date	Arrival_Time_Hour
0	3	0	5	24	3	2019	22	20	22	1
1	1	3	0	1	5	2019	5	50	1	13
2	4	2	1	9	6	2019	9	25	10	4
3	3	3	0	12	5	2019	18	5	12	23
4	3	0	5	1	3	2019	16	50	1	21

```
from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
```

```
data1 = ss.fit_transform(data)
```

```
data1 = pd.DataFrame(data1,columns=data.columns)
data.head()
```

	Airline	Source	Destination	Date	Month	Year	Dep_Time_Hour	Dep_Time_Mins	Arrival_date	Arrival_Time_Hour
0	3	0	5	24	3	2019	22	20	22	1
1	1	3	0	1	5	2019	5	50	1	13
2	4	2	1	9	6	2019	9	25	10	4
3	3	3	0	12	5	2019	18	5	12	23
4	3	0	5	1	3	2019	16	50	1	21

```
y = data1['Price']
x = data1.drop(columns=['Price'],axis=1)
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
x_train.head()
```

	Airline	Source	Destination	Date	Month	Year	Dep_Time_Hour	Dep_Time_Mins	Arrival_date	Arriv
10004	0.864716	0.040721	-0.29563	1.591104	0.250153	0.0	-0.781129	0.297937	1.546321	
3684	0.014369	0.040721	-0.29563	-0.531796	0.250153	0.0	-0.259258	0.297937	-0.461621	
1034	1.715063	0.040721	-0.29563	1.237288	-0.608777	0.0	0.436570	1.097240	1.191978	
3909	0.864716	0.040721	-0.29563	0.883471	-1.467707	0.0	-0.085301	1.363674	0.955750	
3088	-1.261152	0.040721	-0.29563	1.237288	1.109082	0.0	0.784483	-0.501367	1.310092	

```
x_train.shape
```

```
(8545, 11)
```

Double-click (or enter) to edit

```
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, AdaBoostRegressor
rfr=RandomForestRegressor()
gb=GradientBoostingRegressor()
ad=AdaBoostRegressor()
```

```
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

```
for i in [rfr,gb,ad]:
    i.fit(x_train,y_train)
    y_pred=i.predict(x_test)
    test_score=r2_score(y_test,y_pred)
    train_score=r2_score(y_train,i.predict(x_train))
```

```

if abs(train_score-test_score)<=0.2:
    print(i)

    print("R2 score is",r2_score(y_test,y_pred))
    print("R2 for train data",r2_score(y_train,i.predict(x_train)))
    print("Mean Absolute Error is",mean_absolute_error(y_pred,y_test))
    print("Mean Squared Error is",mean_squared_error(y_pred,y_test))
    print("Root Mean Squared Error is",(mean_squared_error(y_pred,y_test,squared=False)))

RandomForestRegressor()
R2 score is 0.850432571284373
R2 for train data 0.9506969959263003
Mean Absolute Error is 0.25475317086441596
Mean Squared Error is 0.14872838541074593
Root Mean Squared Error is 0.38565319318105734
GradientBoostingRegressor()
R2 score is 0.7652955778741217
R2 for train data 0.7338510043179753
Mean Absolute Error is 0.36494013808679193
Mean Squared Error is 0.23338777734765506
Root Mean Squared Error is 0.48310224316148137
AdaBoostRegressor()
R2 score is 0.2533535201196103
R2 for train data 0.25941367300272333
Mean Absolute Error is 0.7258662199205039
Mean Squared Error is 0.7424579427407443
Root Mean Squared Error is 0.8616599925380917

```

```

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.layers import Dense, Activation, Dropout
from tensorflow.keras.optimizers import Adam

```

```

model=keras.Sequential()
model.add(Dense(7,activation='relu',input_dim=11))

```

```

model.add(Dense(7,activation='relu'))

```

```

model.add(Dense(1,activation='linear'))

```

```

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 7)	84
dense_1 (Dense)	(None, 7)	56
dense_2 (Dense)	(None, 1)	8

```

Total params: 148
Trainable params: 148
Non-trainable params: 0

```

```

model.compile(loss='mse', optimizer='rmsprop',metrics=['mae'])

```

```

model.fit(x_train, y_train, batch_size=20, epochs=10)

```

```

Epoch 1/10
428/428 [=====] - 1s 2ms/step - loss: 1.0145 - mae: 0.7881
Epoch 2/10
428/428 [=====] - 1s 2ms/step - loss: 0.8853 - mae: 0.7353
Epoch 3/10
428/428 [=====] - 1s 2ms/step - loss: 0.8238 - mae: 0.7080
Epoch 4/10
428/428 [=====] - 1s 2ms/step - loss: 0.7768 - mae: 0.6834
Epoch 5/10
428/428 [=====] - 1s 2ms/step - loss: 0.7311 - mae: 0.6565
Epoch 6/10

```

```

428/428 [=====] - 1s 2ms/step - loss: 0.6987 - mae: 0.6386
Epoch 7/10
428/428 [=====] - 1s 2ms/step - loss: 0.6760 - mae: 0.6236
Epoch 8/10
428/428 [=====] - 1s 2ms/step - loss: 0.6555 - mae: 0.6091
Epoch 9/10
428/428 [=====] - 1s 2ms/step - loss: 0.6388 - mae: 0.5984
Epoch 10/10
428/428 [=====] - 1s 2ms/step - loss: 0.6243 - mae: 0.5880
<keras.callbacks.History at 0x7f72b828f790>

```

```

from sklearn.model_selection import cross_val_score
for i in range(2,5):
    cv=cross_val_score(rfr,x,y,cv=i)
    print(rfr,cv.mean())

```

```

RandomForestRegressor() 0.7892661446182386
RandomForestRegressor() 0.7926319036646371
RandomForestRegressor() 0.8009278254254942

```

```
from sklearn.model_selection import RandomizedSearchCV
```


```

param_grid={'n_estimators':[10,30,50,70,100], 'max_depth':[None,1,2,3],
            'max_features':['auto','sqrt']}
rfr=RandomForestRegressor()
rf_res=RandomizedSearchCV(estimator=rfr,param_distributions=param_grid,cv=3,verbose=2,n_jobs=-1)

rf_res.fit(x_train,y_train)

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits




```

gb=GradientBoostingRegressor()
gb_res=RandomizedSearchCV(estimator=gb,param_distributions=param_grid,cv=3,verbose=2,n_jobs=-1)

gb_res.fit(x_train,y_train)

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits



```

rfr=RandomForestRegressor(n_estimators=10,max_features='sqrt',max_depth=None)
rfr.fit(x_train,y_train)
y_train_pred=rfr.predict(x_train)
y_test_pred=rfr.predict(x_test)
print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))

```

```

train accuracy 0.9223637224760787
test accuracy 0.7637242809880902

```

```

from sklearn.model_selection import cross_val_score
for i in range(2,5):
    cv=cross_val_score(gb,x,y,cv=i)
    print(rfr,cv.mean())

```

```

RandomForestRegressor(max_features='sqrt', n_estimators=10) 0.7261930738745277
RandomForestRegressor(max_features='sqrt', n_estimators=10) 0.7290564575182709
RandomForestRegressor(max_features='sqrt', n_estimators=10) 0.727707269447425

```

```
gb=GradientBoostingRegressor(n_estimators=10,max_features='sqrt',max_depth=None)
gb.fit(x_train,y_train)
y_train_pred=gb.predict(x_train)
y_test_pred=gb.predict(x_test)
print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))
```

```
train accuracy 0.6364865173543215
test accuracy 0.2501677267836613
```

```
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor

from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

```
knn=KNeighborsRegressor()
svr=SVR()
dt=DecisionTreeRegressor()

for i in [knn,svr,dt]:
    i.fit(x_train,y_train)
    y_pred=i.predict(x_test)
    test_score=r2_score(y_test,y_pred)
    train_score=r2_score(y_train,i.predict(x_train))
    if abs(train_score-test_score)<=0.1:
        print(i)
        print("R2 score is",r2_score(y_test,y_pred))
        print("R2 score for train data",r2_score(y_train,i.predict(x_train)))
        print("Mean Absolute Error is",mean_absolute_error(y_test,y_pred))
        print("Mean Squared Error is",mean_squared_error(y_test,y_pred))
        print("Root Mean Squared Error is",(mean_squared_error(y_test,y_pred,squared=False)))
```

```
KNeighborsRegressor()
R2 score is 0.7337698733752689
R2 score for train data 0.7878038246704271
Mean Absolute Error is 0.35917333406078
Mean Squared Error is 0.26473662896136735
Root Mean Squared Error is 0.5145256348923417
SVR()
R2 score is 0.6248128034087579
R2 score for train data 0.5957444387377182
Mean Absolute Error is 0.4162888458787714
Mean Squared Error is 0.3730824715981053
Root Mean Squared Error is 0.6108047737191526
```

```
knn=KNeighborsRegressor(n_neighbors=2,algorithm='auto',metric_params=None,n_jobs=-1)
knn.fit(x_train,y_train)
y_train_pred=knn.predict(x_train)
y_test_pred=knn.predict(x_test)

print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))
```

```
train accuracy 0.8797607060998262
test accuracy 0.7013502693959782
```

```
from sklearn.model_selection import cross_val_score
for i in range(2,5):
    cv=cross_val_score(knn,x,y,cv=i)
    print(knn,cv.mean())
```

```
KNeighborsRegressor(n_jobs=-1, n_neighbors=2) 0.6301907440142309
KNeighborsRegressor(n_jobs=-1, n_neighbors=2) 0.6458609920294707
KNeighborsRegressor(n_jobs=-1, n_neighbors=2) 0.6646580886084823
```

```
predicted_values=pd.DataFrame({'Actual':y_test,'Predicted':y_pred})
```


predicted_values

	Actual	Predicted
6075	1.641563	1.681688
3544	-0.895161	-0.895161
9290	0.021842	-0.217169
5032	-1.133955	-1.190563
2483	0.826714	1.516636
...
9796	-0.364002	0.976150
9870	-0.968253	-0.614942
10062	-0.354459	-0.636414
8802	-0.439479	-0.466156
8617	1.007382	0.614706

2137 rows × 2 columns

```
prices=rfr.predict(x_test)
```

```
price_list=pd.DataFrame({'Price':prices})
```

price_list

	Price
0	1.354620
1	-0.342400
2	0.036547
3	-1.179241
4	1.243284
...	...
2132	0.842135
2133	-0.825865
2134	-0.508992
2135	0.314771
2136	0.590268

2137 rows × 1 columns

```
import pickle
pickle.dump(rfr,open('model1.pkl','wb'))
```

